



t-academy

Future Skills Academy
for Emerging Technologies

CERTIFIED

DEEP LEARNING

PROFESSIONAL

(CDLP)

ICTC International Council
for Technology Certifications

Digital Tech Faculty Expert (DTeX)



Introduction to CNNs



Deep Learning

Deep Learning

- Welcome to the section on Convolutional Neural Networks!'
- CNNs are a specific architecture of Neural Networks that are extremely effective at dealing with image data.
- Let's review what we will learn in this section!

Deep Learning

CNN Section

- Understanding CNNs
 - Image Kernels and Filters
 - Convolutions
 - Pooling Layers
- MNIST Dataset (Grayscale Images)
- CIFAR-10 Dataset (Color Images)


Deep Learning

CNN Section

- Working with image files (jpg,png,etc...) with CNNs.
- CNNs for malaria blood cell classification.
- CNN exercise on Fashion Image Dataset.

MNIST Data

Deep Learning



Deep Learning


- A classic data set in Deep Learning is the MNIST data set.
- Let's quickly cover some basics about it since we'll be using it quite frequently during this section of the course!

Deep Learning

- Fortunately this data is easy to access with PyTorch Vision. We will download:
 - 60,000 training images
 - 10,000 test images

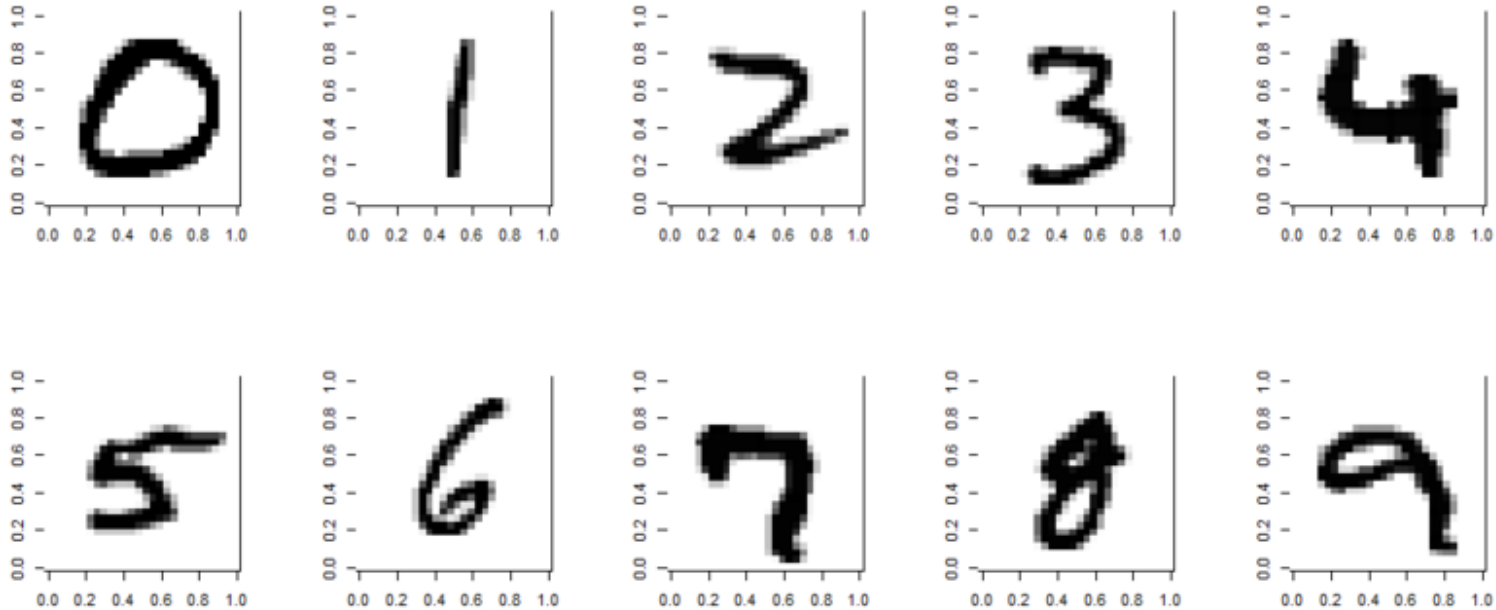
Deep Learning

- The MNIST data set contains handwritten single digits from 0 to 9




Deep Learning

- A single digit image can be represented as an array



Deep Learning

- Specifically, 28 by 28 pixels



- The values represent the grayscale image




Deep Learning

- We will first explore how we could approach this data set with a standard ANN.
- In this case, to feed it into our network we will need to flatten the 28 by 28 array to 784

Deep Learning

Deep Learning


0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	.6	.8	0	0	0	0	0	0
0	0	0	0	0	0	.7	1	0	0	0	0	0	0
0	0	0	0	0	0	.7	1	0	0	0	0	0	0
0	0	0	0	0	0	.5	1	.4	0	0	0	0	0
0	0	0	0	0	0	0	1	.4	0	0	0	0	0
0	0	0	0	0	0	0	1	.4	0	0	0	0	0
0	0	0	0	0	0	0	1	.7	0	0	0	0	0
0	0	0	0	0	0	0	1	1	0	0	0	0	0
0	0	0	0	0	0	0	.9	1	.1	0	0	0	0
0	0	0	0	0	0	0	.3	1	.1	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0



0 0 0 0 0 0 .6 .8 ... 1 .1 0 0 0 0

Deep Learning

0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	.6	.8	0	0	0	0	0	0
0	0	0	0	0	0	.7	1	0	0	0	0	0	0
0	0	0	0	0	0	.7	1	0	0	0	0	0	0
0	0	0	0	0	0	.5	1	.4	0	0	0	0	0
0	0	0	0	0	0	0	1	.4	0	0	0	0	0
0	0	0	0	0	0	0	1	.4	0	0	0	0	0
0	0	0	0	0	0	0	1	.7	0	0	0	0	0
0	0	0	0	0	0	0	1	1	0	0	0	0	0
0	0	0	0	0	0	0	.9	1	.1	0	0	0	0
0	0	0	0	0	0	0	.3	1	.1	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0




0 0 0 0 0 0 .6 .8 ... 1 .1 0 0 0 0


784
inputs

Deep Learning

0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	.6	.8	0	0	0	0	0	0
0	0	0	0	0	0	.7	1	0	0	0	0	0	0
0	0	0	0	0	0	.7	1	0	0	0	0	0	0
0	0	0	0	0	0	.5	1	.4	0	0	0	0	0
0	0	0	0	0	0	0	1	.4	0	0	0	0	0
0	0	0	0	0	0	0	1	.4	0	0	0	0	0
0	0	0	0	0	0	0	1	.7	0	0	0	0	0
0	0	0	0	0	0	0	1	1	0	0	0	0	0
0	0	0	0	0	0	0	.9	1	.1	0	0	0	0
0	0	0	0	0	0	0	.3	1	.1	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0




0 0 0 0 0 0 .6 .8 ... 1 .1 0 0 0 0




Deep Learning

0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	.6	.8	0	0	0	0	0	0
0	0	0	0	0	0	.7	1	0	0	0	0	0	0
0	0	0	0	0	0	.7	1	0	0	0	0	0	0
0	0	0	0	0	0	.5	1	.4	0	0	0	0	0
0	0	0	0	0	0	0	1	.4	0	0	0	0	0
0	0	0	0	0	0	0	1	.4	0	0	0	0	0
0	0	0	0	0	0	0	1	.7	0	0	0	0	0
0	0	0	0	0	0	0	1	1	0	0	0	0	0
0	0	0	0	0	0	0	.9	1	.1	0	0	0	0
0	0	0	0	0	0	0	.3	1	.1	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0




0 0 0 0 0 0 .6 .8 ... 1 .1 0 0 0 0




784
neurons

Deep Learning


0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	.6	.8	0	0	0	0	0	0
0	0	0	0	0	0	.7	1	0	0	0	0	0	0
0	0	0	0	0	0	.7	1	0	0	0	0	0	0
0	0	0	0	0	0	.5	1	.4	0	0	0	0	0
0	0	0	0	0	0	0	1	.4	0	0	0	0	0
0	0	0	0	0	0	0	1	.4	0	0	0	0	0
0	0	0	0	0	0	0	1	.7	0	0	0	0	0
0	0	0	0	0	0	0	1	1	0	0	0	0	0
0	0	0	0	0	0	0	.9	1	.1	0	0	0	0
0	0	0	0	0	0	0	.3	1	.1	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0



0 0 0 0 0 .6 .8 ... 1 .1 0 0 0 0



784
neurons



Hidden
Layers



10 output
neurons

Deep Learning

- Flattening out the image ends up removing some of the 2-D information, such as the relationship of a pixel to its neighboring pixels.
- For now, we'll ignore this, but come back to it later when we discuss CNN in depth!

MNIST ANN

PART ONE: DATA

MNIST ANN

PART TWO: NEURAL NETWORK

MNIST ANN

PART THREE: TRAINING AND EVALUATION

MNIST ANN

PART FOUR: EVALUATION

Image Filters



Deep Learning

Deep Learning

- Before we dive into CNNs, let's first discuss a few key ideas in computer vision.
- Computer vision is a general term of using computer programs to process image data.

Deep Learning

- If you've ever used photo editing software, you have probably seen filters, such as a blur filter. But how do these work?



Deep Learning

- Filters are essentially an **image kernel**, which is a small matrix applied to an entire image.
- Certain popular filters are well known, for example a blur filter:

$$\begin{pmatrix} 0.0625 & 0.125 & 0.0625 \\ 0.125 & 0.25 & 0.125 \\ 0.0625 & 0.125 & 0.0625 \end{pmatrix}$$

Deep Learning

- Let's explore how these image kernel/filters actually get applied to an image.

Deep Learning

- Filters allow us to transform images

0	0	0	0	0	0
0	1	1	1	1	0
0	1	-1	-1	1	0
0	1	-1	-1	1	0
0	1	1	1	1	0
0	0	0	0	0	0

Deep Learning

- Here we have an grayscale image

0	0	0	0	0	0
0	1	1	1	1	0
0	1	-1	-1	1	0
0	1	-1	-1	1	0
0	1	1	1	1	0
0	0	0	0	0	0

Deep Learning

- Values scaled between -1 and 1

0	0	0	0	0	0
0	1	1	1	1	0
0	1	-1	-1	1	0
0	1	-1	-1	1	0
0	1	1	1	1	0
0	0	0	0	0	0

Deep Learning

- Filters allow us to transform images


0	0	0	0	0	0
0	1	1	1	1	0
0	1	-1	-1	1	0
0	1	-1	-1	1	0
0	1	1	1	1	0
0	0	0	0	0	0

0	0	0
0	-1	1
0	1	-1

3 by 3 Filter

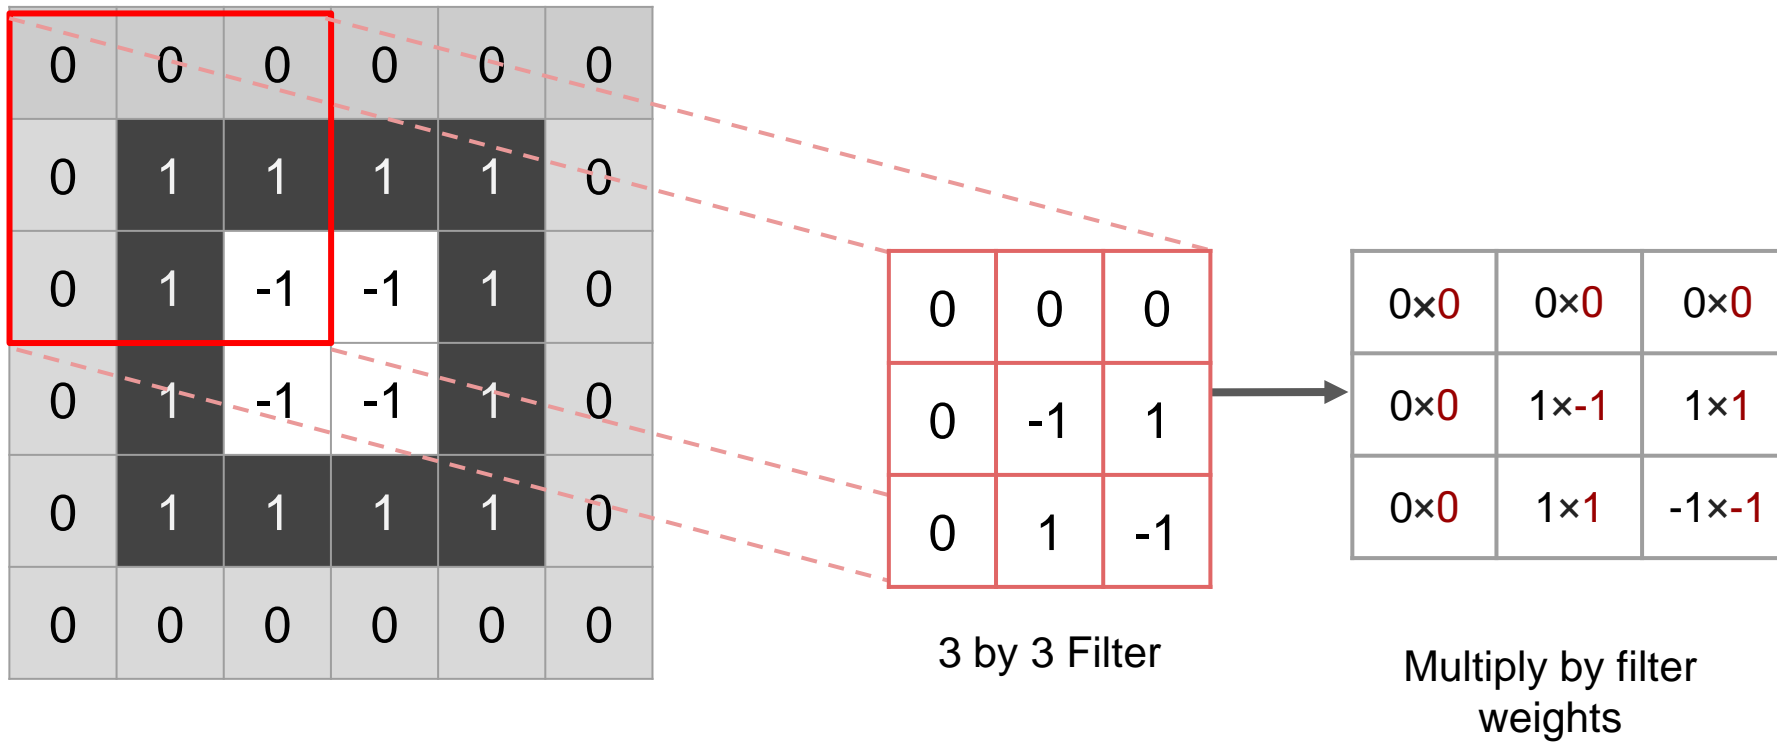
Deep Learning

- Filters allow us to transform images



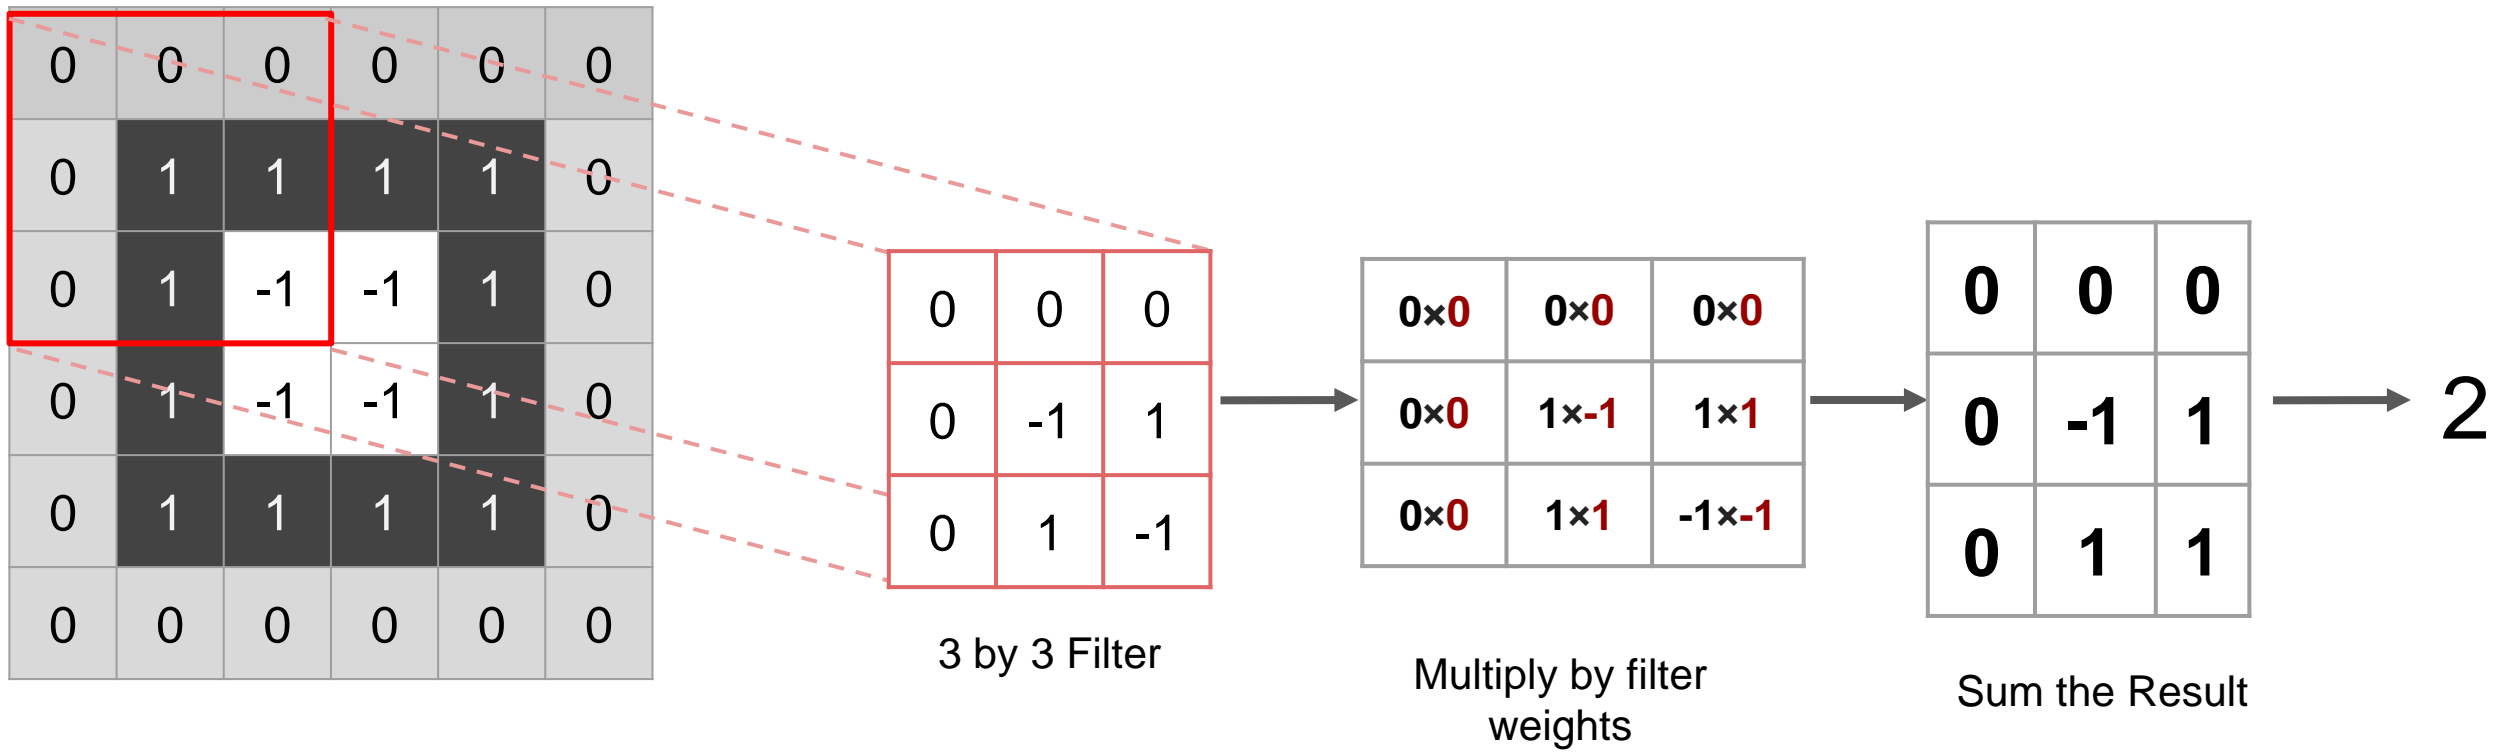
Deep Learning

- Filters are essentially a matrix




Deep Learning

- Notice how the resolution will decrease




Deep Learning

- We can also edit our stride distance



Deep Learning

- Stride Distance of 2 Example



Deep Learning

Let's explore an interactive example!


- **setosa.io/ev/image-kernels/**

Deep Learning

- In the context of CNNs, these “filters” are referred to as **convolution kernels**.
- The process of passing them over an image is known as **convolution**.
- Let's go over a few more important factors!


Deep Learning

During convolution, we would lose borders



Deep Learning


During convolution, we would lose borders



3 by 3 Filter


Deep Learning

We can **pad** the image with more values



Deep Learning

This allows us to preserve the image size



Deep Learning

- Now that we understand image filters, let's explore the architecture of a CNN that allows the network to come up with the best weights for a filter in the **convolutional layer**.

Convolutional Layers

Deep Learning



Deep Learning

- Recall that running an ANN for the MNIST data set resulted in a network with relatively good accuracy.
- However, there are some issues with always using ANN models for image data.

Deep Learning

ANNs

- Large amount of parameters (over 100,000 for tiny 28 by 28 images)
- We lose all 2D information by flattening out the image
- Will only work on very similar, well centered images.

Deep Learning

- A CNN can use convolutional layers to help alleviate these issues.
- A convolutional layer is created when we apply multiple image filters to the input images.
- The layer will then be trained to figure out the best filter weight values.

Deep Learning


- A CNN also helps reduce parameters by focusing on **local connectivity**.
- Not all neurons will be fully connected.
- Instead, neurons are only connected to a subset of local neurons in the next layer (these end up being the filters!)

Deep Learning

- Let's understand this local connectivity and its connection to filters by beginning with a simplified 1-D example.
- We will later expand this to 2-D inputs for a grayscale image and then later on to 3-D tensor inputs for color images.


Deep Learning

- ANN




Deep Learning

- Fully Connected, lots of parameters!




Deep Learning

- CNN -
Convolutional
Layer




Deep Learning

- CNN -
Localized
Connections




Deep Learning

- CNN - Here we have 1 filter



Deep Learning


- CNN - Here we have 2 filters



Deep Learning


- We just saw a 1-D example of convolution, but recall, grayscale images are 2-D!
- Plus we want to preserve that 2-D relational information in the convolutional layer.

Deep Learning




**Input
Image**

Deep Learning




**Input
Image**

Deep Learning




**Input
Image**

Deep Learning




**Input
Image**

Deep Learning




**Input
Image**

Deep Learning




**Input
Image**

Deep Learning




**Input
Image**

Deep Learning




**Input
Image**


Deep Learning




**Input
Image**




Deep Learning



**Input
Image**



Deep Learning



**Input
Image**


**Convolutio
nal
Layer**

Deep Learning

- But this was just in 2D for grayscale images.
- What about color images?
- Color Images can be thought of as 3D Tensors consisting of Red, Green, and Blue color channels.


Deep Learning

- Additive color mixing allows us to represent a wide variety of colors by simply combining different amounts of R, G, B.




Deep Learning

- RGB allows to produce a range of colors




Deep Learning

- Each color channel will have intensity values.
- You may have already seen this sort of representation in other software with RGB sliders.



Deep Learning



- The shape of the color array then has 3 dimensions.
- Height
- Width
- Color Channels


Deep Learning

This means when you read in an image and check its shape, it will look something like:


- **(1280,720,3)**
- **1280** pixel width
- **720** pixel height
- **3** color channels

Deep Learning

- This means when you read in an image and check its shape, it will look something like:
 - **(720,1280,3)**
 - **720** pixel height
 - **1280** pixel width
 - **3** color channels



Deep Learning




Keep in mind the computer won't "know" a channel is Red, it just knows that there are now 3 intensity channels.


Deep Learning

- How do we then perform a convolution on a color image?
- We end up with a 3D filter, with values for each color channel.

Deep Learning




Input Layer




Filter

Deep Learning



Input Layer



Filter 1

Filter 2 ...

Deep Learning


- Often convolutional layers are fed into another convolutional layer.
- This allows the networks to discover patterns within patterns, usually with more complexity for later convolutional layers.

Deep Learning

- We've learned a lot so far and we have one final theory topic to cover before coding our own CNNs, and that is the Pooling Layer! (also known as a downsampling layer)


Deep Learning

- ANN



Pooling Layers

Deep Learning




Deep Learning

- Even with local connectivity, when dealing with color images and possibly 10s or 100s of filters we will have a large amount of parameters.
- We can use pooling layers to reduce this.


Deep Learning

- Pooling layers accept convolutional layers as input:




Deep Learning

- Our convolutional layers will often have many filters




Deep Learning

- Our convolutional layers will often have many filters



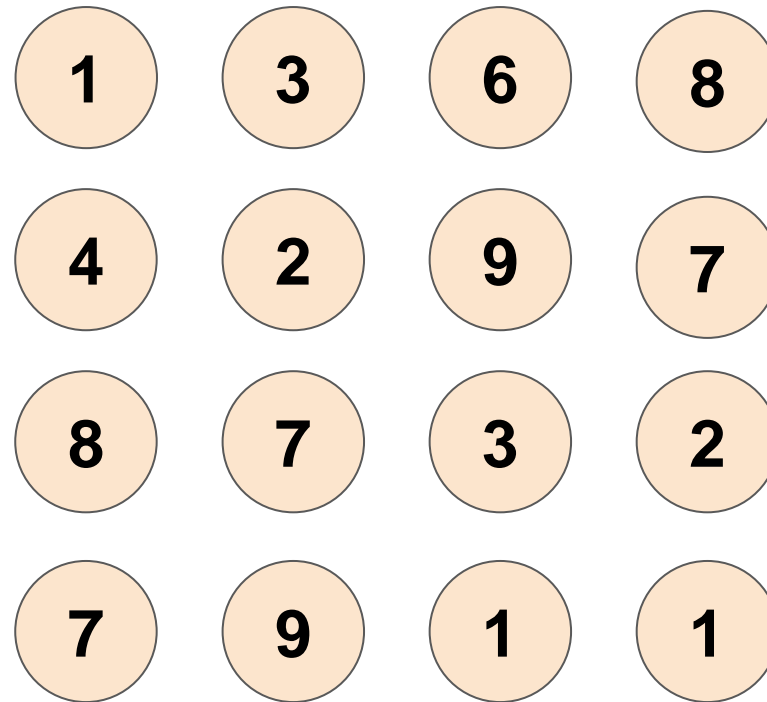
Deep Learning

- Our convolutional layers will often have many filters



Deep Learning


- Let's imagine a filter



1	3	6	8
4	2	9	7
8	7	3	2
7	9	1	1

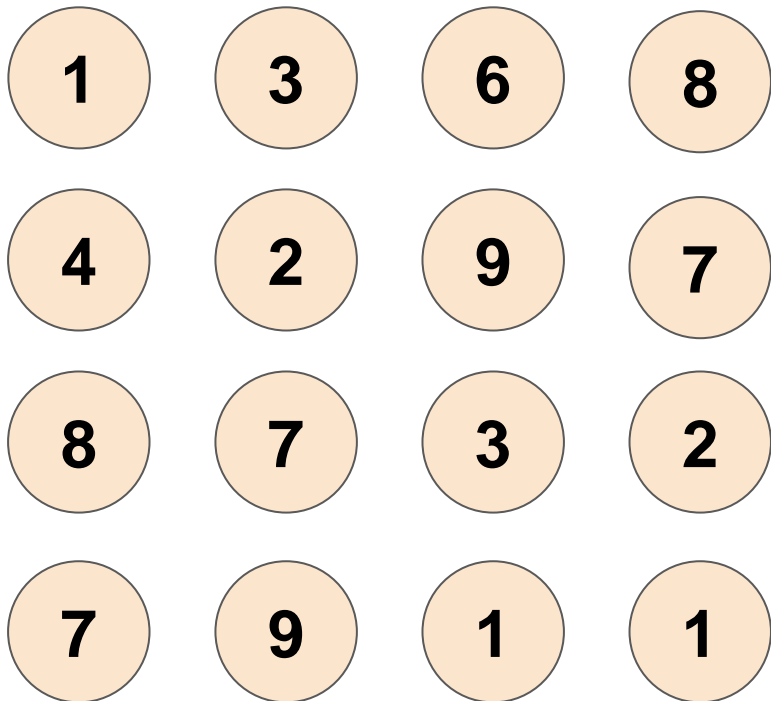
Deep Learning

- We can reduce the size with subsampling



Deep Learning


- We can reduce the size with subsampling



1	3	6	8
4	2	9	7
8	7	3	2
7	9	1	1



Deep Learning

- Max Pooling: Window 2x2 , Stride: 2




Deep Learning

- Max Pooling: Window 2x2 , Stride: 2




Deep Learning

- Max Pooling: Window 2x2 , Stride: 2





Deep Learning

- Max Pooling: Window 2x2 , Stride: 2





Deep Learning

- Max Pooling: Window 2x2 , Stride: 2





Deep Learning

- Max Pooling: Window 2x2 , Stride: 2




Deep Learning

- Average Pooling: Window 2x2 , Stride: 2



Deep Learning

- This greatly reduces our number of parameters!



Deep Learning

- This pooling layer will end up removing a lot of information, even a small pooling “kernel” of 2 by 2 with a stride of 2 will remove 75% of the input data.

Deep Learning

- Another common technique deployed with CNN is called “Dropout”
- Dropout can be thought of as a form of regularization to help prevent overfitting.
- During training, units are randomly dropped, along with their connections.

Deep Learning

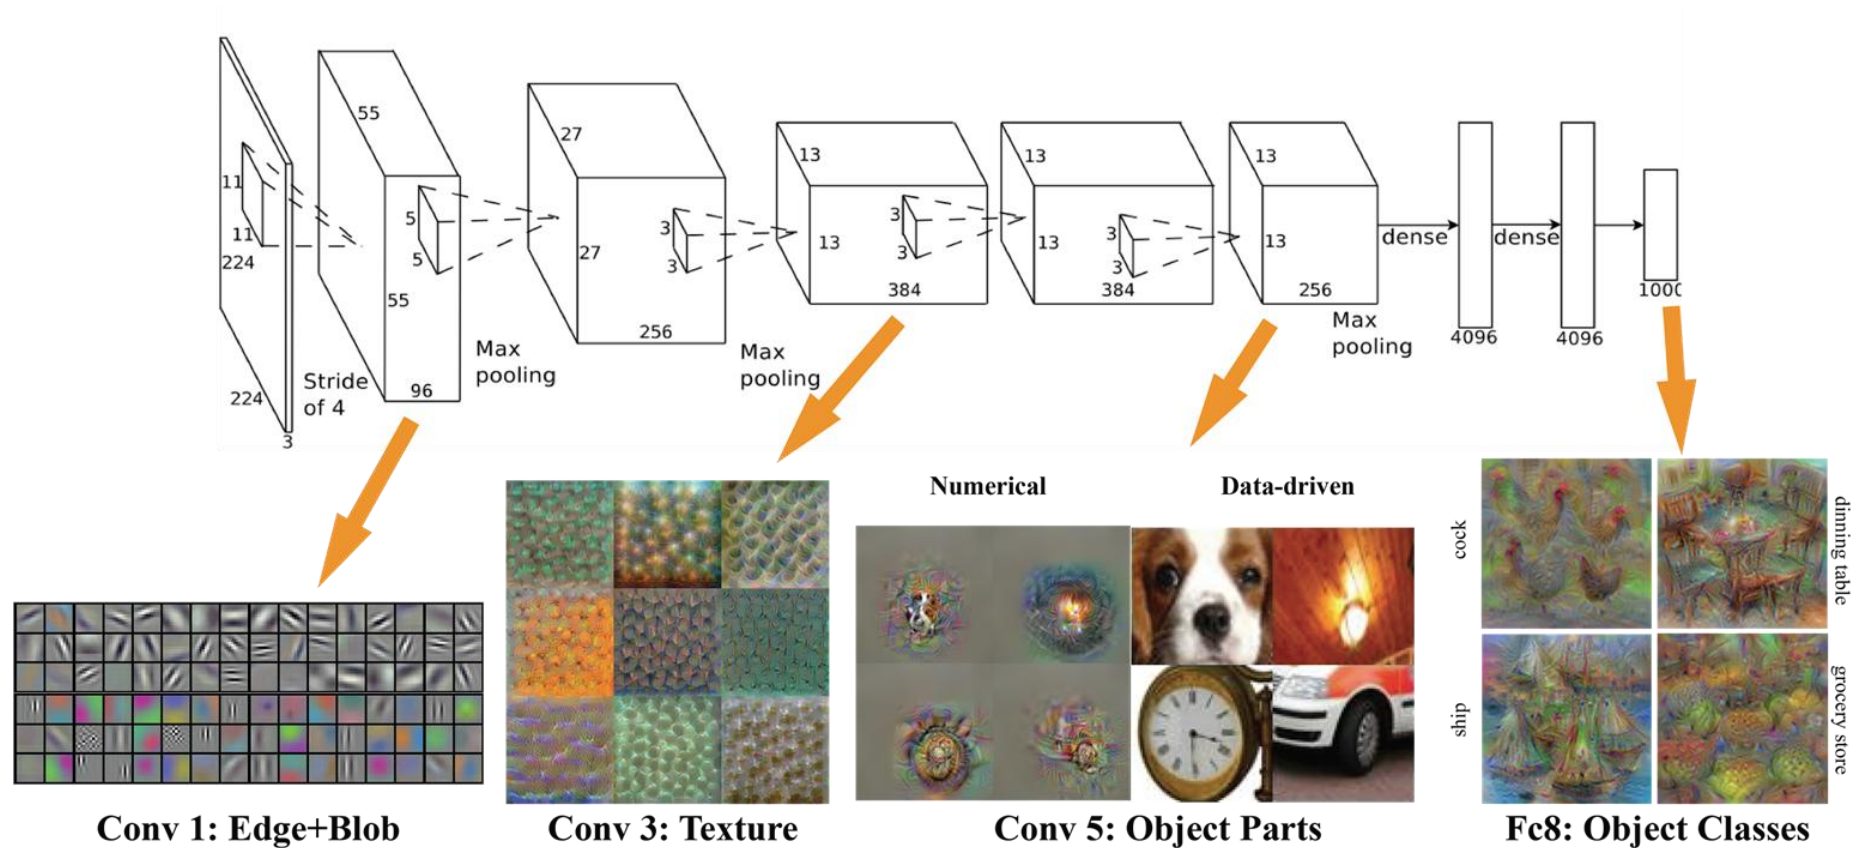
- This helps prevent units from “co-adapting” too much.
- Let’s also quickly point out some famous CNN architectures!

Deep Learning

- LeNet-5 by Yann LeCun
- AlexNet by Alex Krizhevsky et al.
- GoogLeNet by Szegedy at Google Research
- ResNet by Kaiming He et al.
- Check out the resource links to the papers discussing these architectures!


Deep Learning

- AlexNet




Deep Learning

- AlexNet




Deep Learning

- CNNs can have all types of architectures!




Deep Learning

- CNNs can have all types of architectures!




Deep Learning

- CNNs can have all types of architectures!




Deep Learning

- CNNs can have all types of architectures!



MNIST Data Revisited

Deep Learning




Deep Learning

- Recall that flattening out the MNIST data caused us to lose 2D information.
- With CNNs, we can feed in the data as an array of 2D images.

Deep Learning

- We can think of the entire group of the 60,000 images as a tensor (an n-dimensional array)



Deep Learning

- For the labels we'll use One-Hot Encoding.
- This means that instead of having labels such as “One”, “Two”, etc... we'll have a single array for each image.

Deep Learning

- The label is represented based off the index position in the label array.
- The corresponding label will be a 1 at the index location and zero everywhere else.
- For example, 4 would have this label array:
 - `[0,0,0,0,1,0,0,0,0,0]`

Deep Learning

Keep in mind that when dealing with tensors of image data, we actually end up with 4 dimensions:

- Number of Images
- Height
- Width
- Color Channels

MNIST with CNN

PART ONE - The Data

MNIST with CNN

PART TWO - Model and Training

MNIST with CNN

PART THREE - Model Evaluation

CIFAR-10 with CNN

PART ONE - Data

Deep Learning

CIFAR-10 dataset are 32by32 images of 10 different objects:

- Airplane, Car, Bird, Cat, Deer, Dog, Frog, Horse, Ship, Truck
- These are color images!

Deep Learning

- In this program, we will be reusing a lot of the previous CNN code, so we will only focus on new additions due to the introduction of 3 color channels (RGB).

CIFAR-10 with CNN

PART TWO - Model Evaluation

Downloading Image Data

Deep Learning



Deep Learning

- So far we've only dealt with pre-packaged data sets such as MNIST and CIFAR-10.
- But what about working with real image files? Like .jpg or .png?

Deep Learning

- Let's learn about TensorFlow's built in tools for generating image data batches from directories with real files.
- First we need to download a large zip file of data.

Deep Learning

- This dataset is available as a zip file called **cell_images.zip**
- The link is a supplemental resource in this lecture, it is also within the same Google Drive resource as the slides.
- Let's explore getting the data!

CNN on Custom Images

Part 1 - The Data

CNN on Custom Images

Part 2 Image Data Generator

CNN on Custom Images

Part 3 Creating the Model

CNN on Custom Images

Part 4 Model Evaluation



Let's get started!

CNN Exercises

CNN Exercises Solutions