



t-academy

Future Skills Academy  
for Emerging Technologies

**CERTIFIED**

**DEEP LEARNING**

**PROFESSIONAL**

**(CDLP)**


**ICTC** International Council  
for Technology Certifications

**Digital Tech Faculty Expert (DTeX)**



# Recurrent Neural Networks

Deep Learning



# Deep Learning

- We've used Neural Networks to solve Classification and Regression problems, but we still haven't seen how Neural Networks can deal with sequence information.

# Deep Learning

- Just as CNNs were more effective for use with 2D image data, RNNs are more effective for sequence data (e.g. time-stamped sales data, sequence of text, heart beat data, etc...)


# Deep Learning

- Time Series




# Deep Learning

- Time Series



# Deep Learning

- Sequences



# Deep Learning

## Section Overview

- RNN Theory
- LSTMs and GRU Theory
- Basic Implementation of RNN
- Time Series with an RNN
- Exercise and Solution



# Recurrent Neural Networks Theory

Deep Learning



# Deep Learning

## Examples of Sequences

- Time Series Data (Sales)
- Sentences
- Audio
- Car Trajectories
- Music

# Deep Learning


- Let's imagine a sequence:
  - [1,2,3,4,5,6]
- Would you be able to predict a similar sequence shifted one time step into the future?
  - [2,3,4,5,6,7]

# Deep Learning

- To do this properly, we need to somehow let the neuron “know” about its previous history of outputs.
- One easy way to do this is to simply feed its output back into itself as an input!


# Deep Learning

- Normal Neuron in Feed Forward Network



# Deep Learning


- Recurrent Neuron



- Sends output back to itself!
- Let's see what this looks like over time!


# Deep Learning

- Recurrent Neuron



# Deep Learning


- Recurrent Neuron






# Deep Learning

- Recurrent Neuron



# Deep Learning

- Recurrent Neuron



# Deep Learning


- Cells that are a function of inputs from previous time steps are also known as *memory cells*.
- RNN are also flexible in their inputs and outputs, for both sequences and single vector values.

# Deep Learning

- We can also create entire layers of Recurrent Neurons...


# Deep Learning

- ANN Layer with 3 Neurons:




# Deep Learning

- RNN Layer with 3 Neurons:



# Deep Learning

- “Unrolled” layer.




# Deep Learning

- RNN are also very flexible in their inputs and outputs.
- Let's see a few examples.




# Deep Learning

- Sequence to Sequence (Many to Many)




# Deep Learning

- Given 5 previous words, predict the next 5




# Deep Learning

- Sequence to Vector (Many to One)




# Deep Learning

- Given 5 previous words, predict next word




# Deep Learning

- Vector to Sequence (One to Many)



# Deep Learning

- Given 1 word predict the next 5 words



# Deep Learning

- A basic RNN has a major disadvantage, we only really “remember” the previous output.
- It would be great if we could keep track of longer history, not just short term history.

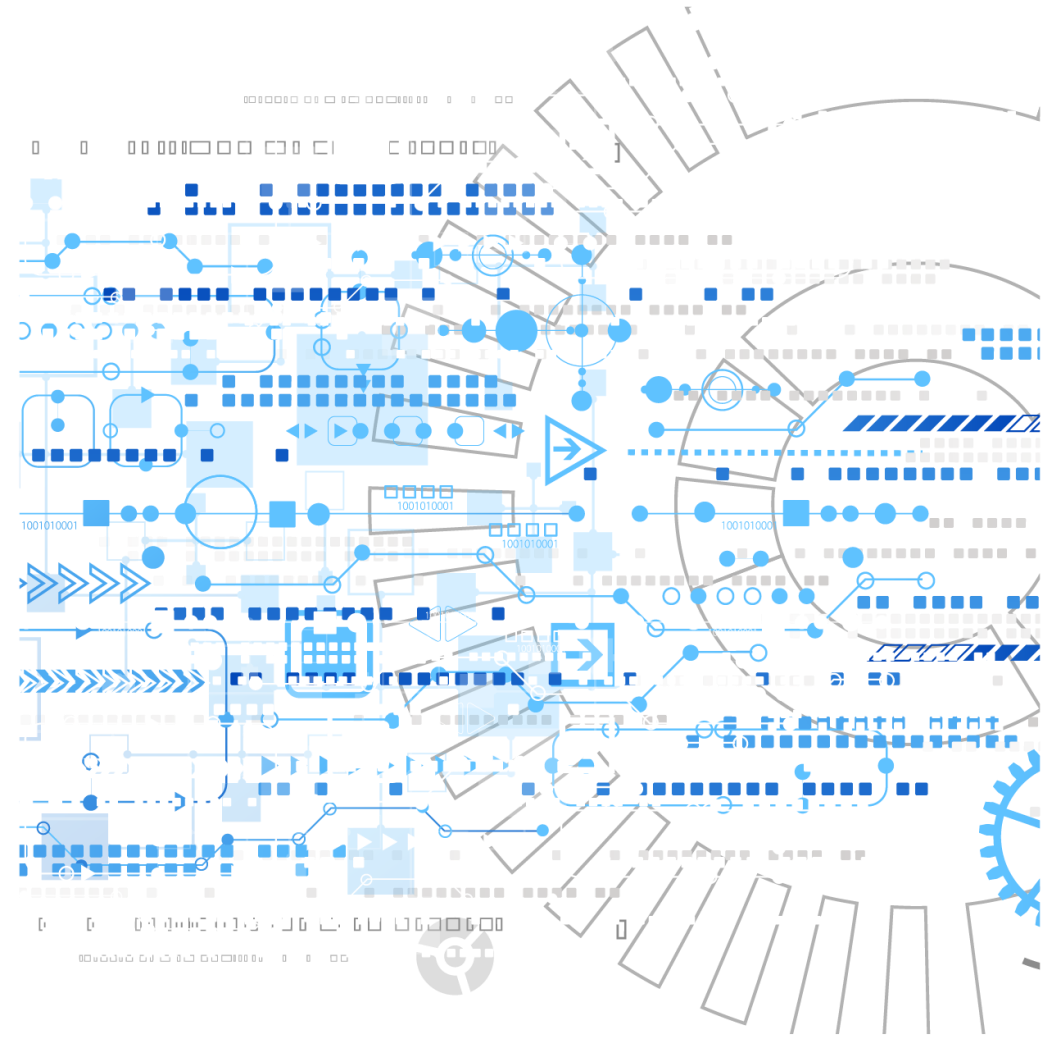
# Deep Learning

- Another issue that arises during training is the “vanishing gradient”.
- Let’s explore vanishing gradients in more detail before moving on to discussing LSTM (Long Short Term Memory Units).



# Exploding and Vanishing Gradients

Deep Learning




# Deep Learning

- As our networks grow deeper and more complex, we have 2 issues arise:
  - Exploding Gradients
  - Vanishing Gradients
- Recall that the gradient is used in our calculation to adjust weights and biases in our network.


# Deep Learning

Let's think about a network.




# Deep Learning

For complex data we need deep networks



# Deep Learning

Issues can arise during backpropagation



# Deep Learning

- Backpropagation goes backwards from the output to the input layer, propagating the error gradient.
- For deeper networks issues can arise from backpropagation, vanishing and exploding gradients!

# Deep Learning

- As you go back to the “lower” layers, gradients often get smaller, eventually causing weights to never change at lower levels.
- The opposite can also occur, gradients explode on the way back, causing issues.


# Deep Learning

- Let's discuss why this might occur and how we can fix it.
- Then in the next lecture we'll discuss how these issues specifically affect RNN and how to use LSTM and GRU to fix them.



# Deep Learning


Why does this happen?



$$f(x) = \frac{1}{1 + e^{-(x)}}$$

# Deep Learning

Why does this happen?




$$f(x) = \frac{1}{1 + e^{-(x)}}$$

$$z = wx + b$$

# Deep Learning

The derivative can be much smaller!




# Deep Learning

- When  $n$  hidden layers use an activation like the sigmoid function,  $n$  small derivatives are multiplied together.
- The gradient could decrease exponentially as we propagate down to the initial layers.

# Deep Learning

## Using Different Activation Functions

Output




$$z = wx + b$$

# Deep Learning

The ReLu doesn't saturate positive values.

Output




$$z = wx + b$$

# Deep Learning

“Leaky” ReLU


Output



$$z = wx + b$$

# Deep Learning

## Exponential Linear Unit (ELU)





# Deep Learning

- Another solution is to perform batch normalization, where your model will normalize each batch using the batch mean and standard deviation.

# Deep Learning

- Choosing different initialization of weights can also help alleviate these issues (Xavier Initialization).

# Deep Learning


- Apart from batch normalization, researchers have also used “gradient clipping”, where gradients are cut off before reaching a predetermined limit (e.g. cut off gradients to be between -1 and 1)

# Deep Learning

- RNN for Time Series present their own gradient challenges, let's explore special LSTM (Long Short Term Memory) neuron units that help fix these issues!

# LTSM and GRU Units

Deep Learning



# Deep Learning

- Many of the solutions previously presented for vanishing gradients can also apply to RNN: different activation functions, batch normalizations, etc...
- However because of the length of time series input, these could slow down training

# Deep Learning

- A possible solution would be to just shorten the time steps used for prediction, but this makes the model worse at predicting longer trends.

# Deep Learning

- Another issue RNN face is that after awhile the network will begin to “forget” the first inputs, as information is lost at each step going through the RNN.
- We need some sort of “long-term memory” for our networks.




# Deep Learning

- The LSTM (Long Short-Term Memory) cell was created to help address these RNN issues.
- Let's go through how an LSTM cell works!


# Deep Learning

A typical RNN




# Deep Learning

RNN




# Deep Learning

RNN




# Deep Learning

RNN



# Deep Learning

RNN




# Deep Learning

LSTM



# Deep Learning

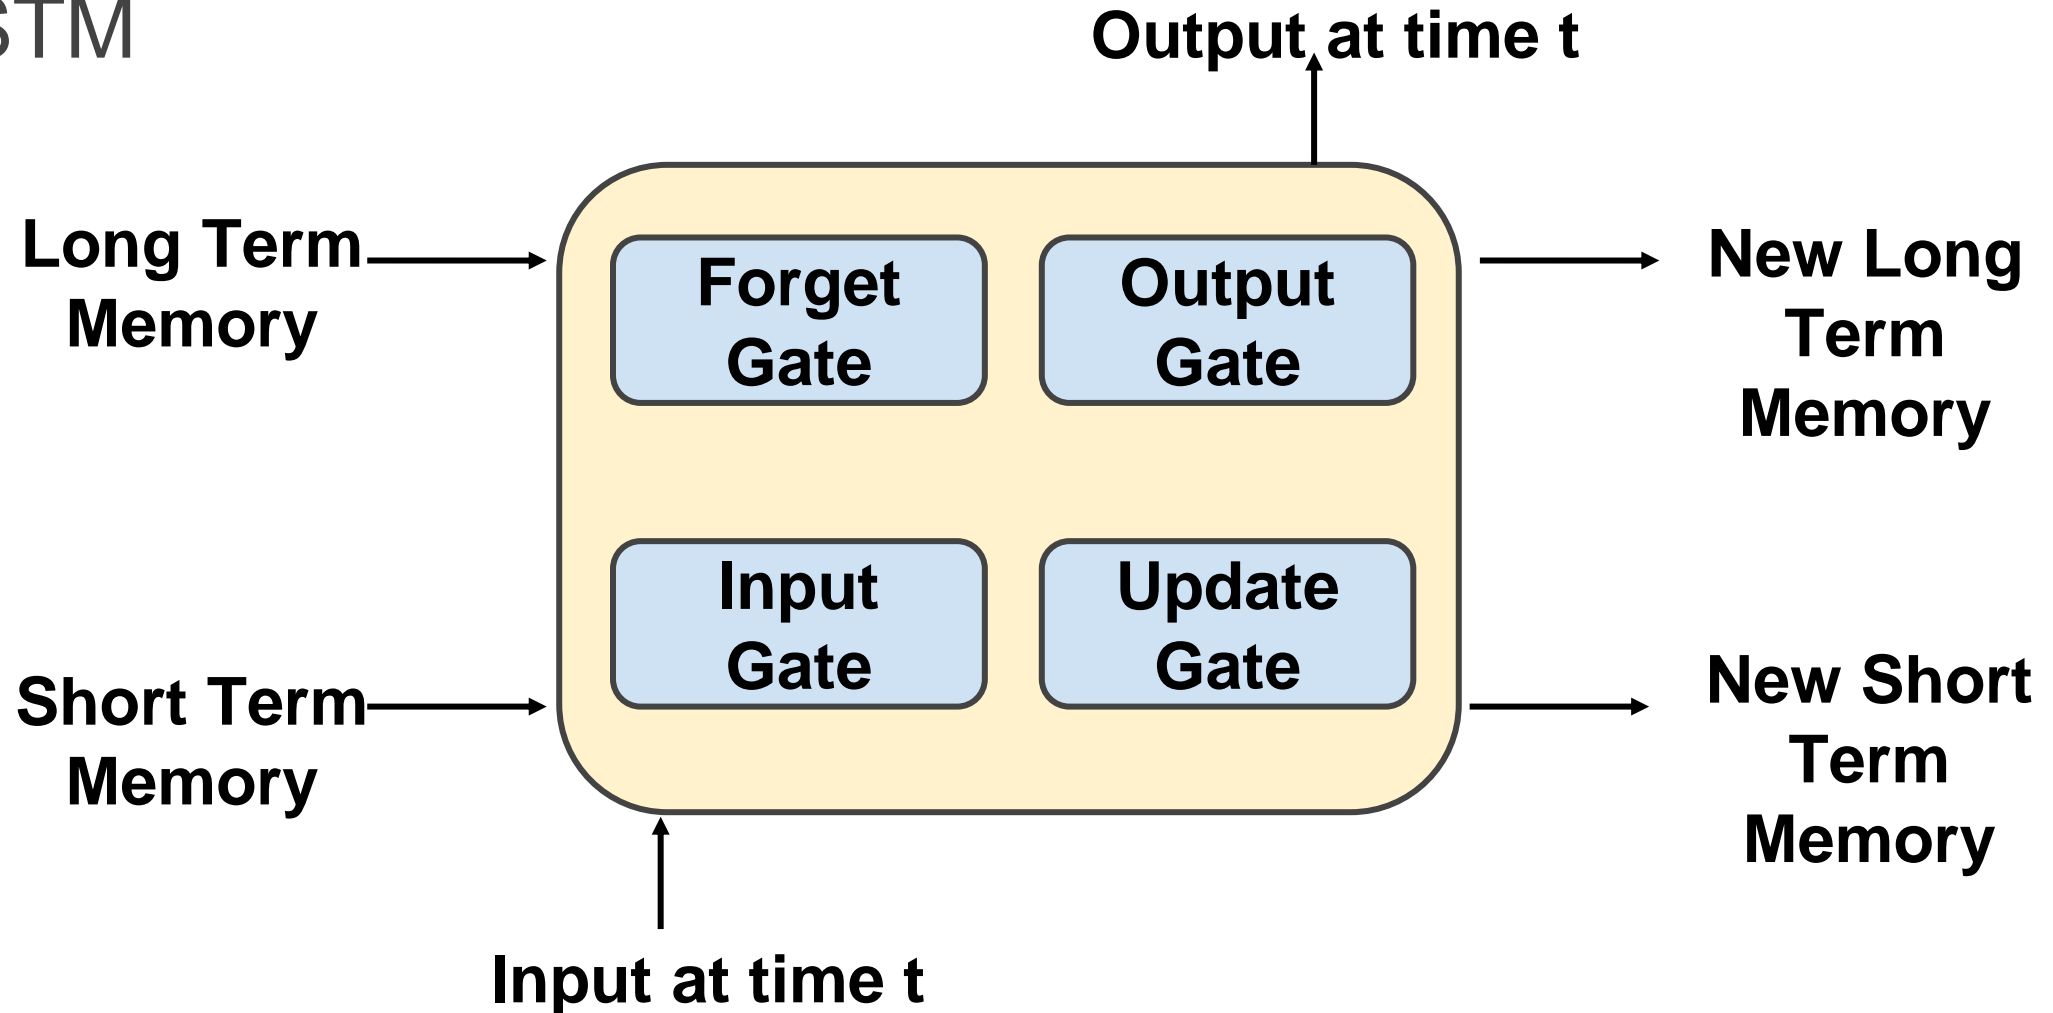
## LSTM






# Deep Learning

## LSTM




# Deep Learning

Gates optionally let information through




# Deep Learning

## LSTM

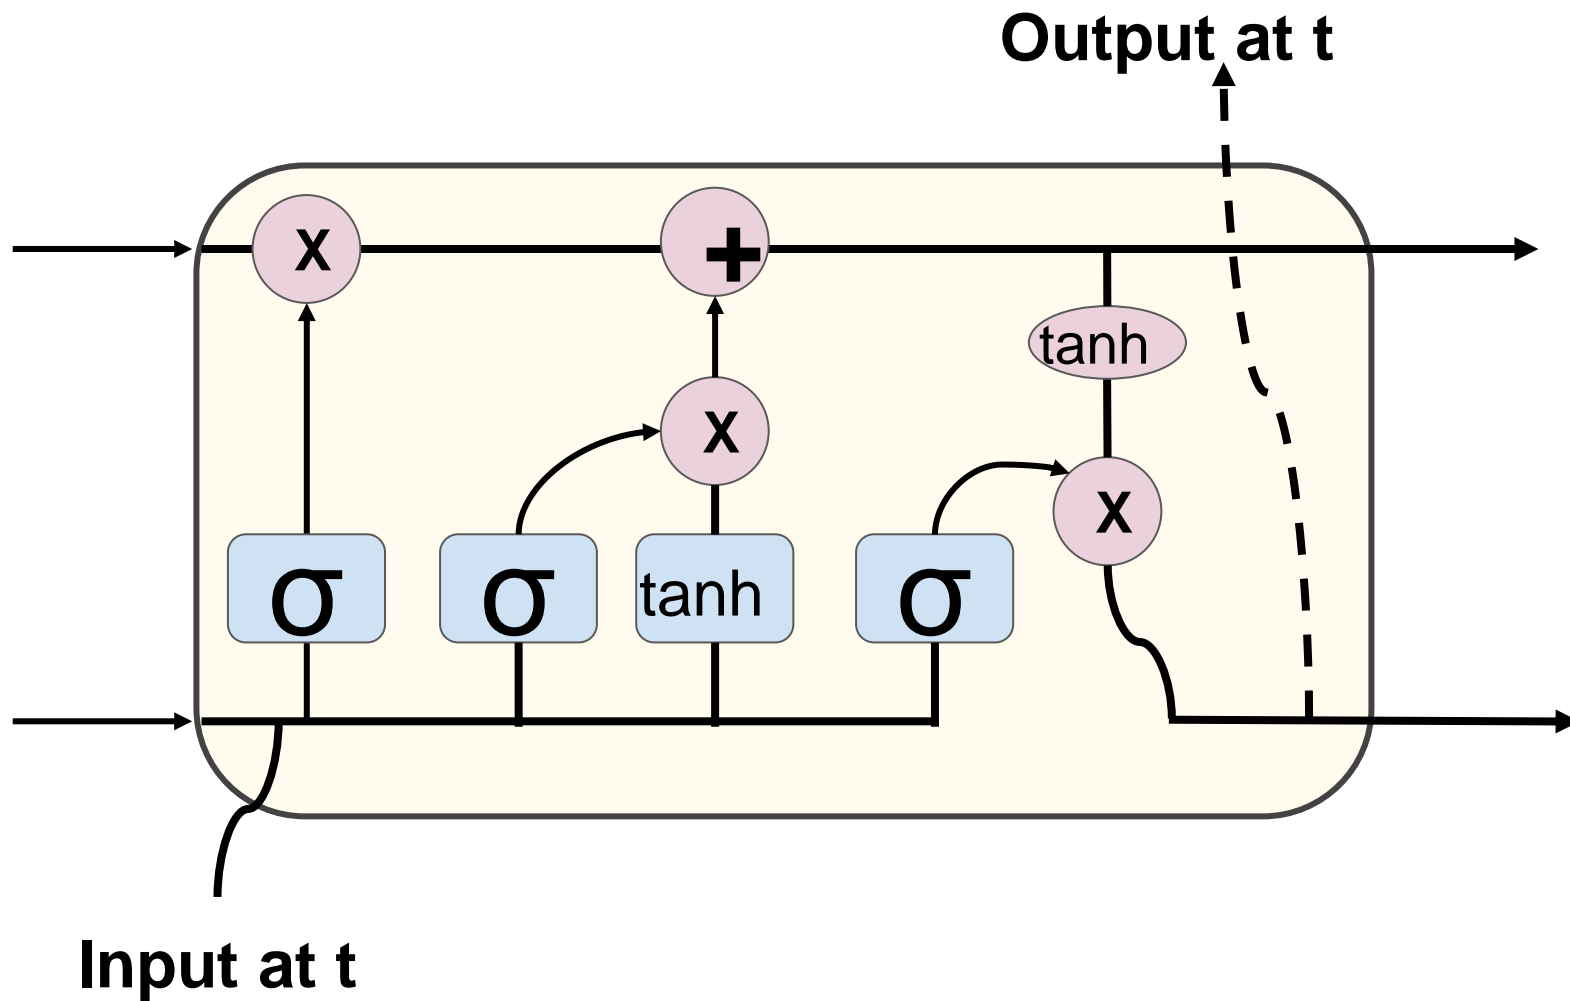


# Deep Learning

## LSTM




# An LSTM Cell



Here we can see the entire LSTM cell.

Let's go through the process!

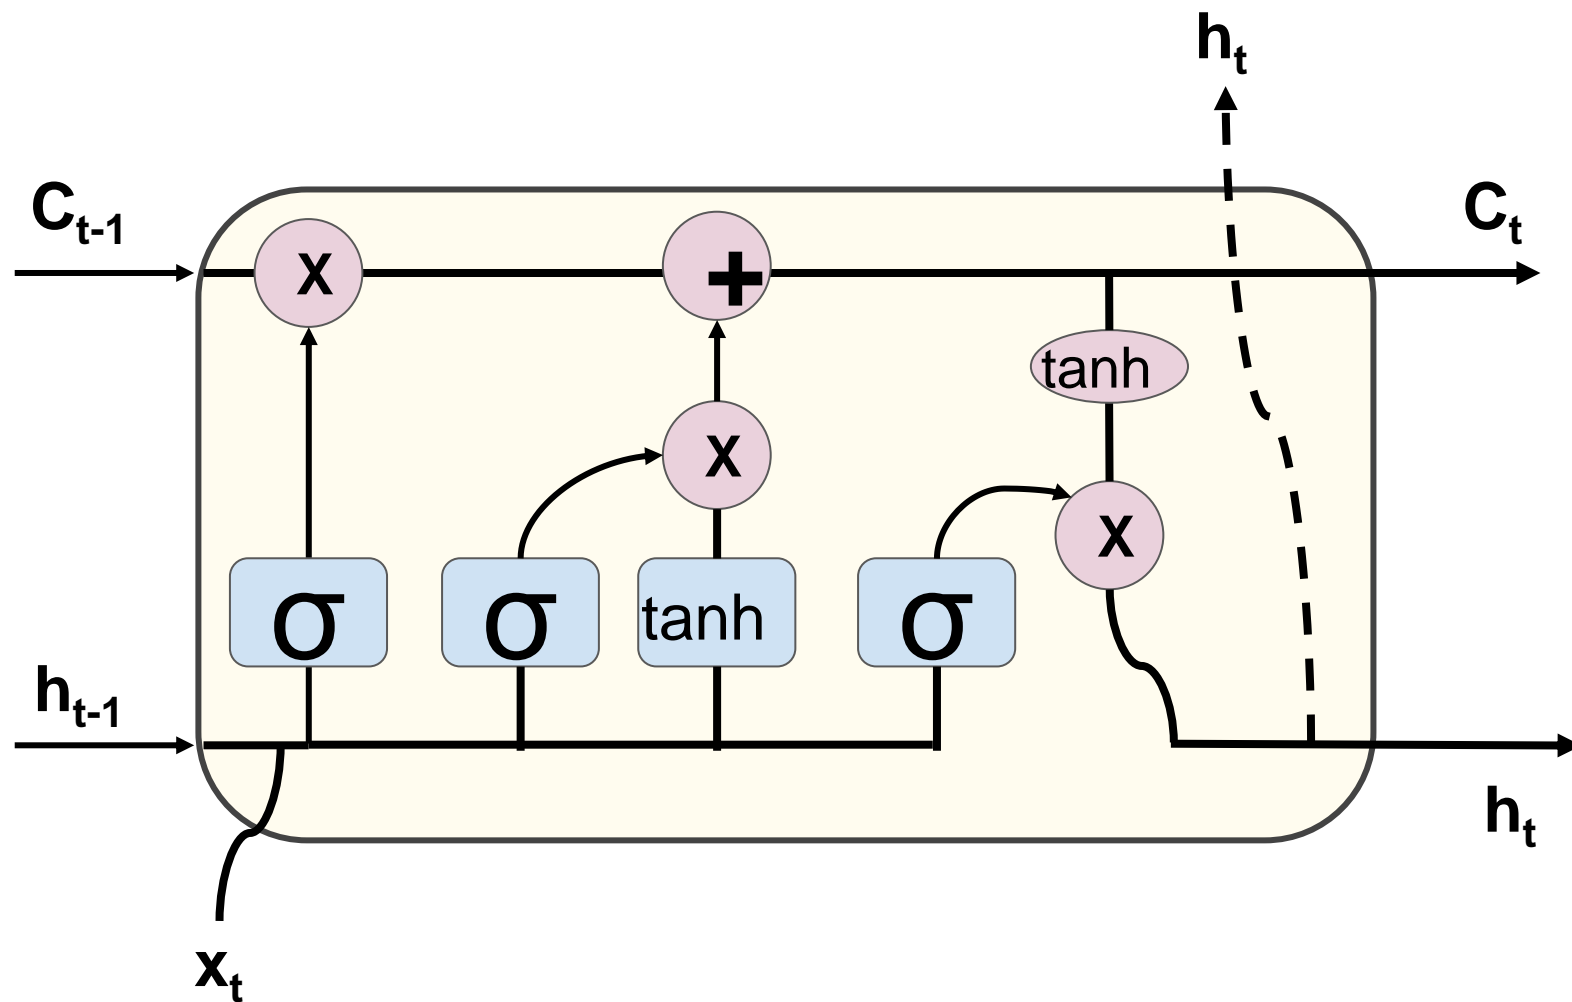
# An LSTM Cell



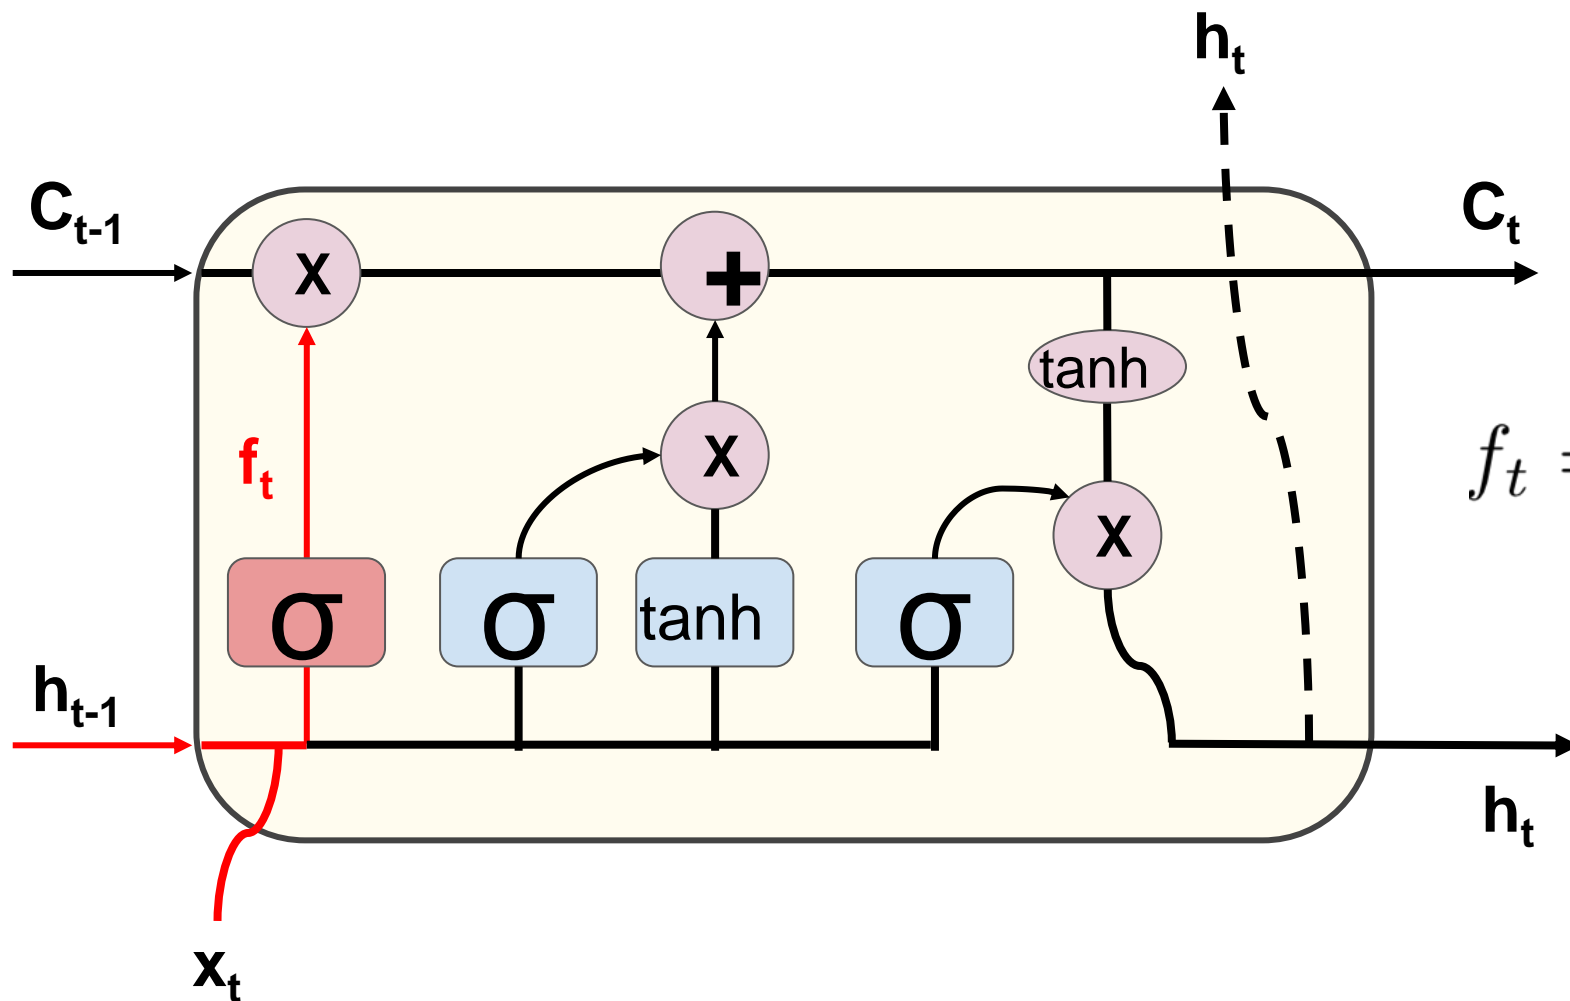
Here we can see the entire LSTM cell.

Let's go through the process!

# An LSTM Cell



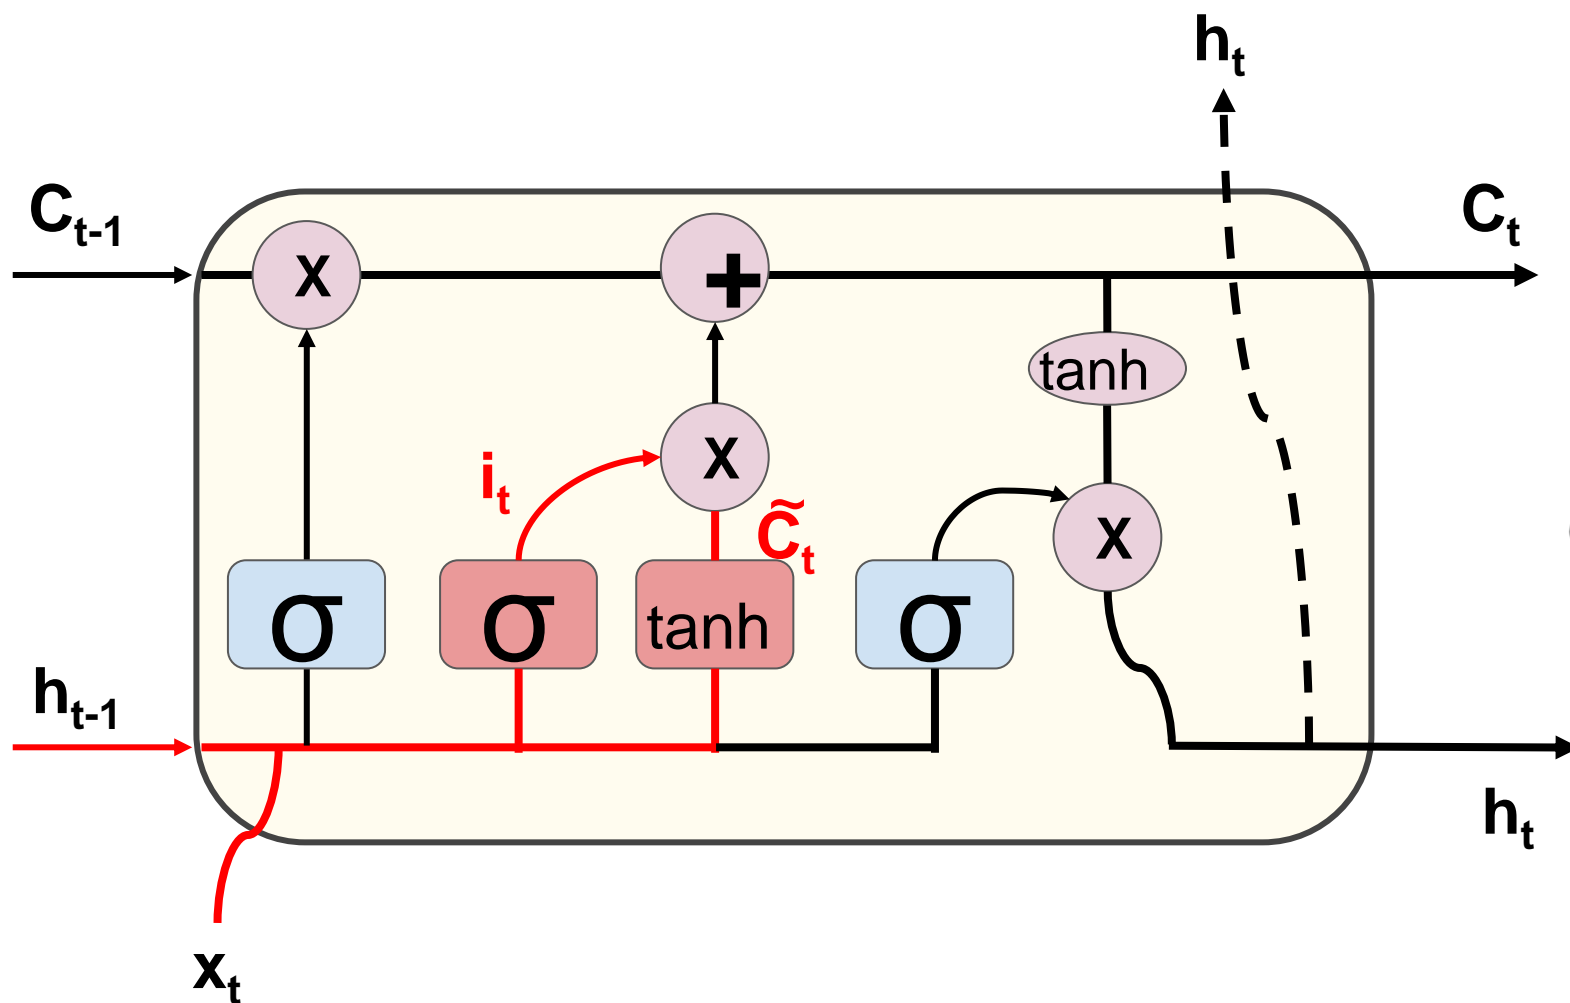
# An LSTM Cell



$$f_t = \sigma (W_f \cdot [h_{t-1}, x_t] + b_f)$$



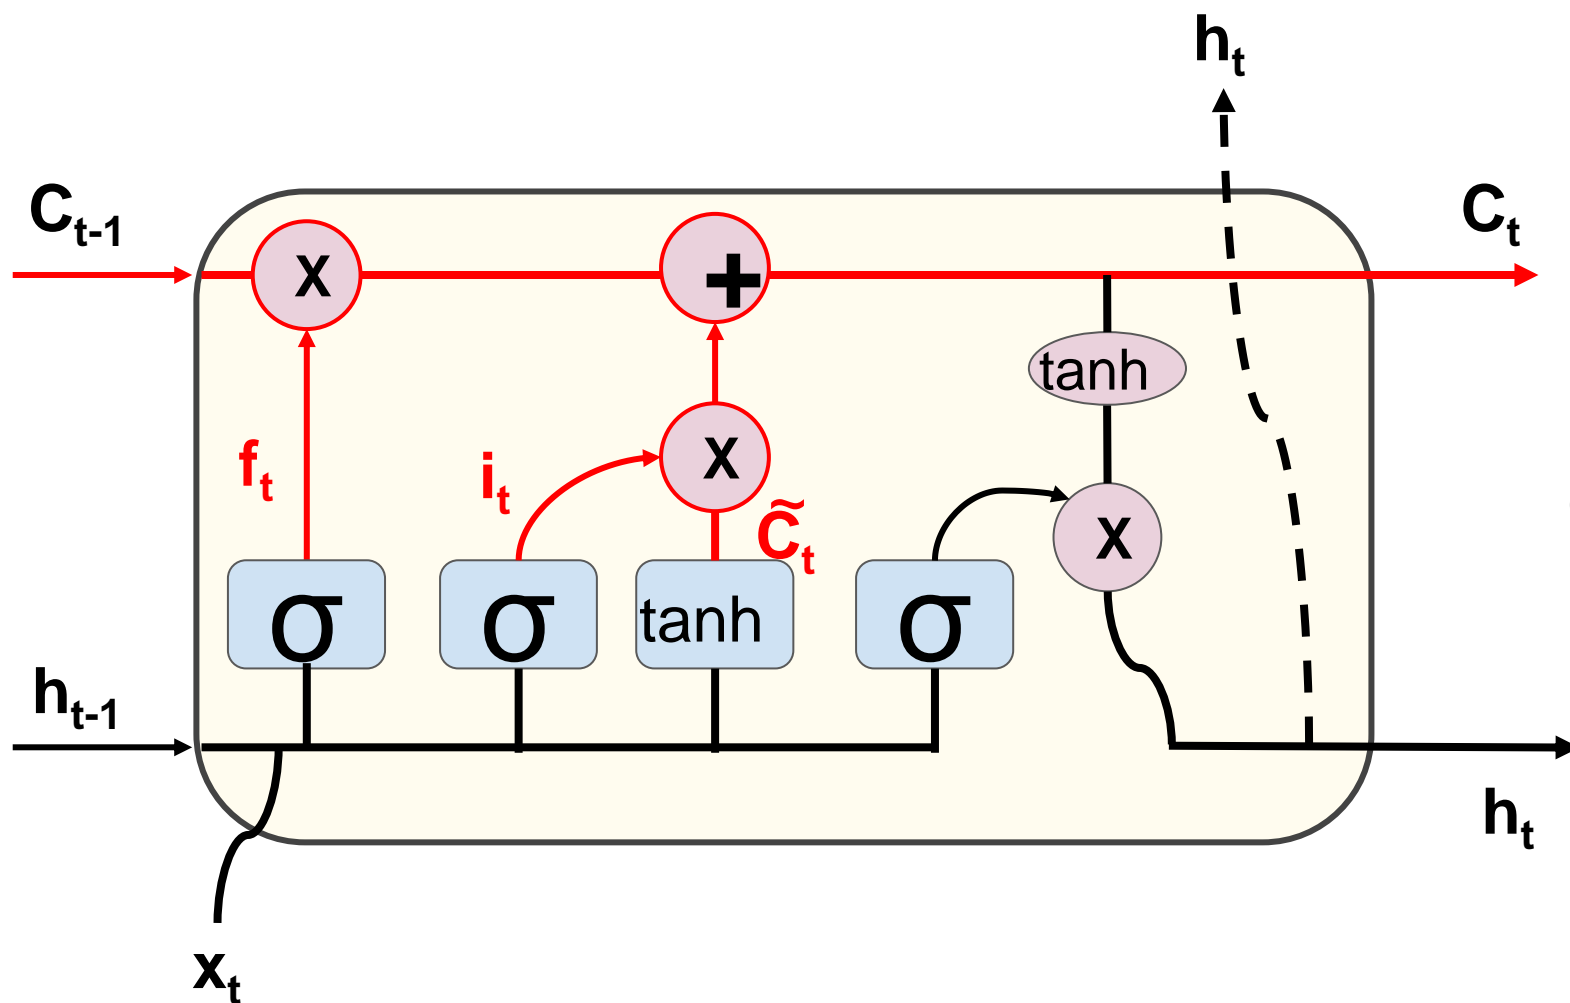
# An LSTM Cell



$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

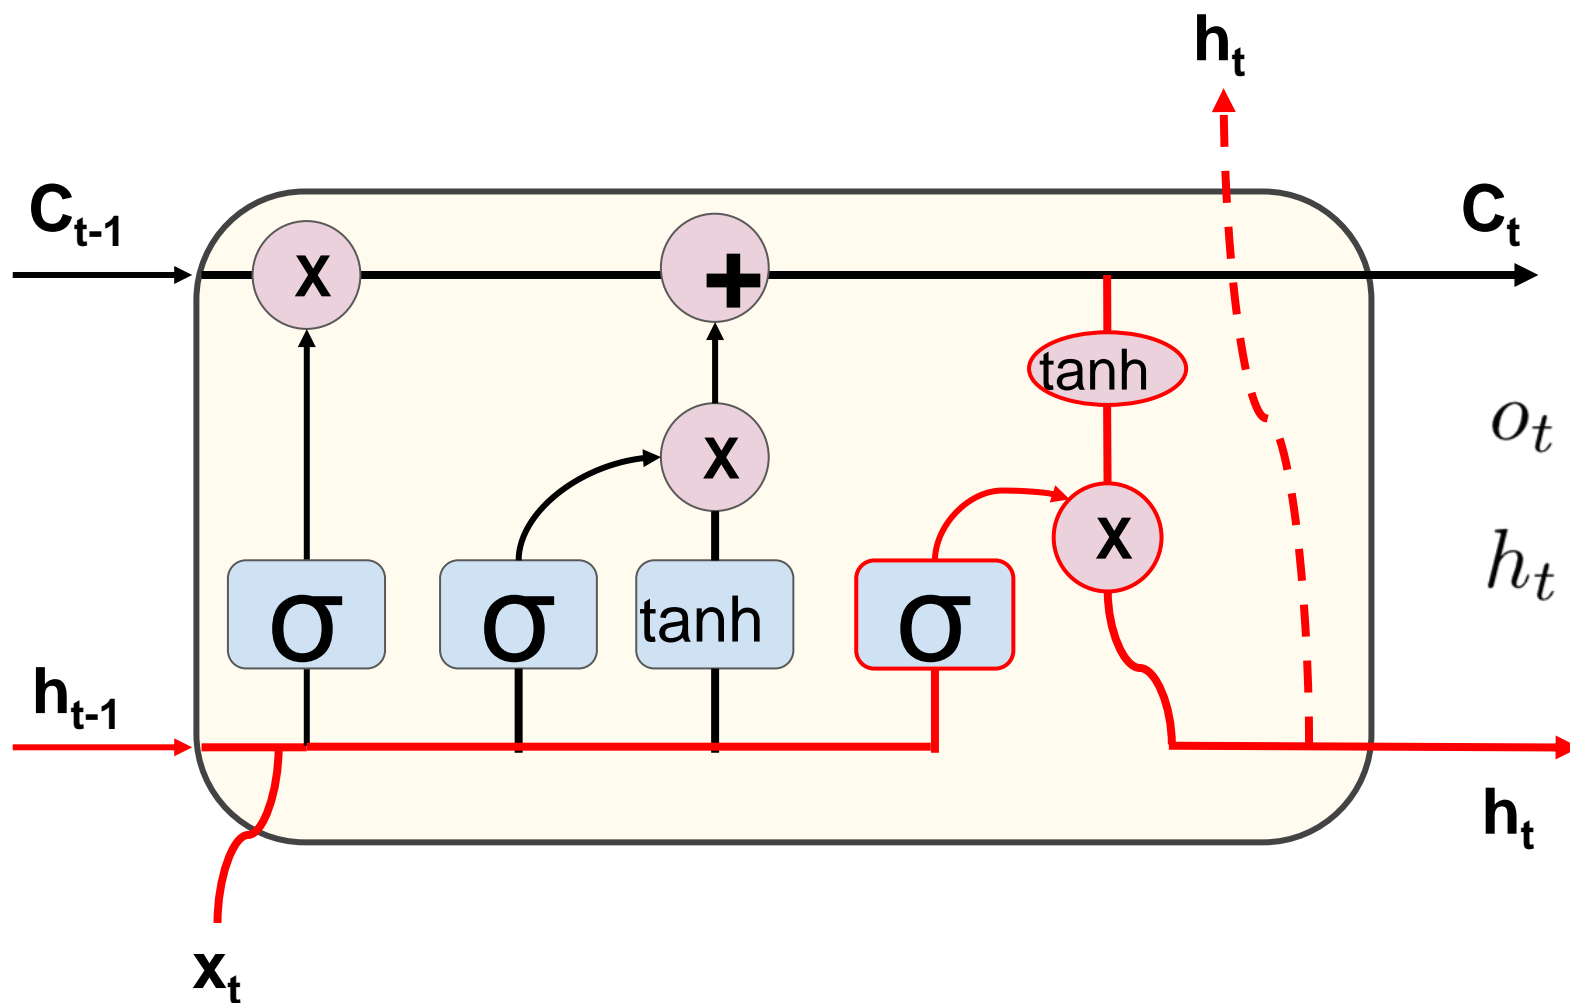
$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

# An LSTM Cell



$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

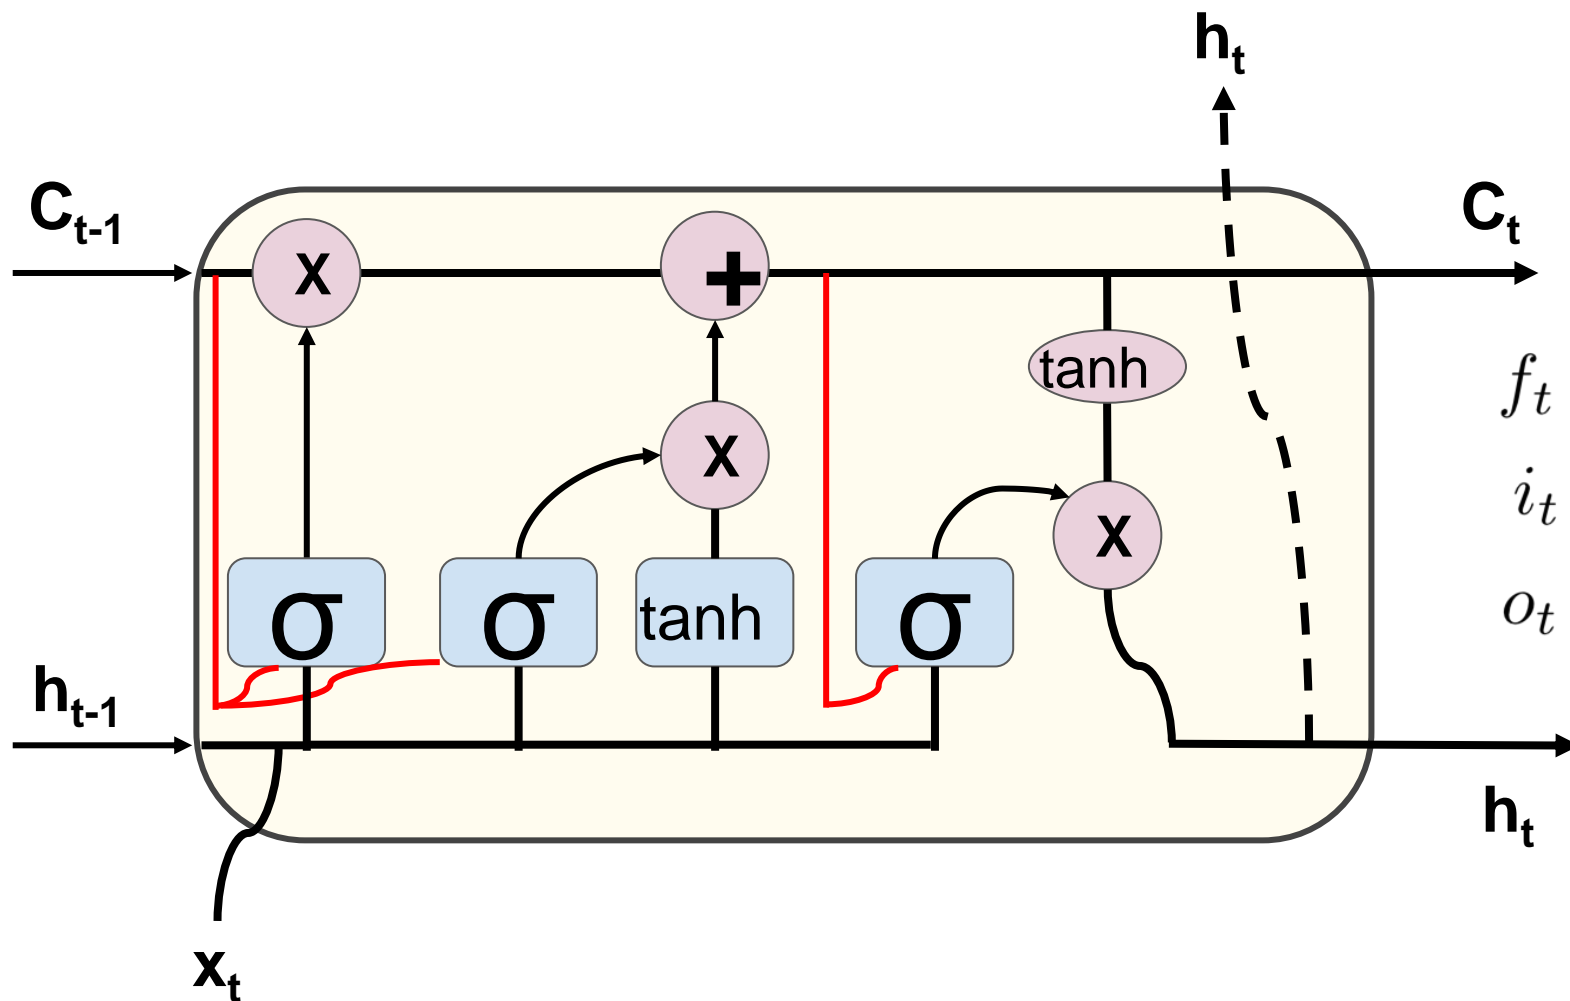
# An LSTM Cell



$$o_t = \sigma(W_o [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh(C_t)$$

# An LSTM Cell with “peepholes”




$$f_t = \sigma(W_f \cdot [C_{t-1}, h_{t-1}, x_t] + b_f)$$

$$i_t = \sigma(W_i \cdot [C_{t-1}, h_{t-1}, x_t] + b_i)$$

$$o_t = \sigma(W_o \cdot [C_t, h_{t-1}, x_t] + b_o)$$

# Gated Recurrent Unit (GRU)



$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma(W_r \cdot [h_{t-1}, x_t])$$


$$\tilde{h}_t = \tanh(W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

# Deep Learning

- Fortunately with our deep learning python library , we simply need to call the import for RNN or LSTM instead of needing to code all of this ourselves!
- Let's explore how to use LSTMs with Python code!

# Basic RNN




Deep Learning

# Deep Learning

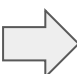
- Let's now explore how to use RNN on a basic time series, such as a sine wave.
- Before we jump to the notebook, let's quickly discuss what RNN sequence batches look like.



# Deep Learning

- Let's imagine a simple time series:
  - [0,1,2,3,4,5,6,7,8,9]
- We separate this into 2 parts:
  - [0,1,2,3,4,5,6,7,8] [9]
- Given **training sequence**,  predict the **next sequence value**.

# Deep Learning

- Keep in mind we can usually decide how long the training sequence and predicted label should be:
  - **[0,1,2,3,4]**  **[5,6,7,8,9]**

# Deep Learning


- We can also edit the size of the training point, as well as how many sequences to feed per batch:
  - **[0,1,2,3]** → **[4]**
  - **[1,2,3,4]** → **[5]**
  - **[2,3,4,5]** → **[6]**

# Deep Learning

- So how do we decide how long the training sequence should be?
- There is no definitive answer, but it should be at least long enough to capture any useful trend information.


# Deep Learning

- For example, if dealing with seasonal data:



# Deep Learning




- If this is monthly, we should include at least 12 months in the training sequence



# Deep Learning

- This often takes domain knowledge and experience, as well as simply experimenting and using RMSE to measure error of forecasted predictions.
- Typically a good starting choice for the label is just one data point into the future.

# Deep Learning

- How do we forecast with RNNs?
  - Let's imagine all our data is:
    - **[0,1,2,3,4,5,6,7,8,9]**
  - And we trained on sequences such as:
    - **[0,1,2,3]** **[4]**
    - **[1,2,3,4]** **[5]** 
    - **[2,3,4,5]** **[6]** 
- 



# Deep Learning

- Then our forecasting technique is to predict a time step ahead, and then incorporate our prediction into the next sequence we predict off of.
- Let's walk through a quick example!

# Deep Learning

- How do we forecast with RNNs?
- Let's imagine all our data is:
  - **[0,1,2,3,4,5,6,7,8,9]**
- And we trained on sequences such as:
  - **[0,1,2,3]** → **[4]**
  - **[1,2,3,4]** → **[5]**
  - **[2,3,4,5]** → **[6]**

# Deep Learning

- **[6,7,8,9]**  **[10]** Forecast prediction!

# Deep Learning

- **[6,7,8,9]**  **[10]** Forecast prediction!
- Then to keep predicting further:

# Deep Learning

- **[6,7,8,9]**  $\Rightarrow$  **[10]** Forecast prediction!

Then to keep predicting further:

- **[7,8,9,10]**  $\Rightarrow$  **[11.2]**

# Deep Learning

- **[6,7,8,9]**  $\Rightarrow$  **[10]** Forecast prediction!

Then to keep predicting further:

- **[7,8,9,10]**  $\Rightarrow$  **[11.2]**
- **[8,9,10,11.2]**  $\Rightarrow$  **[12.4]**

# Deep Learning

- **[6,7,8,9]**  $\Rightarrow$  **[10]** Forecast prediction!

Then to keep predicting further:

- **[7,8,9,10]**  $\Rightarrow$  **[11.2]**
- **[8,9,10,11.2]**  $\Rightarrow$  **[12.4]**
- **[9,10,11.2,12.4]**  $\Rightarrow$  **[14]**

# Deep Learning

- **[6,7,8,9]**  $\Rightarrow$  **[10]** Forecast prediction!

Then to keep predicting further:


- **[7,8,9,10]**  $\Rightarrow$  **[11.2]**
- **[8,9,10,11.2]**  $\Rightarrow$  **[12.4]**
- **[9,10,11.2,12.4]**  $\Rightarrow$  **[14]**
- **[10,11.2,12.4,14]**  $\Rightarrow$  Completed Forecast



# Deep Learning

- Let's explore this further with Python!


# Basic RNN on a Sine Wave



Deep Learning


# Deep Learning

- Let's now train and evaluate our RNN
- Recall our original data:




# Deep Learning

- We split this into **train\_set** and **test\_set**:




# Deep Learning

- During training, we could evaluate performance on the test data.




# Deep Learning

- Let's zoom in on that range!




# Deep Learning

- As we train, we will forecast on this range to visually see the training.



# Deep Learning


- As we train, we will forecast on this range to visually see the training.





# Deep Learning

- As we train, we will forecast on this range to visually see the training.




# Deep Learning

- Eventually once we are satisfied with the results on the test portion, it will be time to forecast into the unknown future.
- This means **retraining** on **all** available data (both train and test) and forecasting beyond the scope of the original data.


# Deep Learning

- Keep in mind, for real data, we would no longer have values to compare these predictions to!



# RNN on a Time Series

Deep Learning




# RNN Exercises

# **RNN Exercises Solutions**

# Multivariate Time Series with LSTM RNNs

Deep Learning



# Deep Learning

- As a quick bonus, we'll discuss how to use LSTMs and RNNs to predict multivariate time series.
- Keep in mind, there are a few cons to using LSTMs for this approach!



# Deep Learning

- As with all neural networks, the model is essentially a black box, difficult to interpret.
- Also there are many well-studied alternatives that are simpler, such as SARIMAX and VARMAX models.

# Deep Learning

- We highly recommend you try those more conventional approaches before settling on LSTMs or RNNs for multivariate time series data.
- Fortunately, setting up for multivariate data only requires 2 main changes.

# Deep Learning

- Multivariate Time Series:
  - Change input shape in LSTM layer to reflect 2-D structure
  - Final dense layer should have a neuron per feature/variable.
- Let's explore these changes!